

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library



Partial evaluation in Python

Partial evaluation in Python

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ What is partial evaluation
- ▶ When should we apply it?
- ▶ Real-world case study
- ▶ Partial evaluation library

Part 1

- **Part 1: What is partial evaluation?**
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

What is partial evaluation?

What is partial evaluation?


- **Part 1: What is partial evaluation?**
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

Partial evaluation is:

- ▶ an optimization technique
- ▶ that generates specialized code
- ▶ using information, available at runtime

Original program


- **Part 1: What is partial evaluation?**
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library



```
>>> power = lambda x, n: x ** n
>>> power(3, 5)
243
```

Separate inputs


- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library



```
v = [power(x, 3) + 3 * power(x, 47)  
    for x in A_LOT_OF_DATA]
```

Specialized program

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library



```
v = [power_3(x) + 3 * power_47(x)  
     for x in A_LOT_OF_DATA]
```

An example: power function

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power(x, n):  
    if not isinstance(n, int) or n < 0:  
        raise ValueError  
    elif n == 0:  
        return 1  
    elif n % 2 == 0:  
        v = power(x, n / 2)  
        return v * v  
    else:  
        return x * power(x, n - 1)
```


Specialized: $n = \text{'foo'}$, $n = 1$

- **Part 1: What is partial evaluation?**
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power_foo(x):  
    raise ValueError  
  
def power_1(x):  
    return x
```

Specialized: $n = 5$

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power_5(x):  
    _pow_2 = x * x  
    _pow_4 = _pow_2 * _pow_2  
    return x * _pow_4
```

Specialized: $n = 27$

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power_27(x):  
    _pow_2 = x * x  
    _pow_3 = _pow_2 * x  
    _pow_6 = _pow_3 * _pow_3  
    _pow_12 = _pow_6 * _pow_6  
    _pow_13 = _pow_12 * x  
    _pow_26 = _pow_13 * _pow_13  
    _pow_27 = _pow_26 * x  
    return _pow_27
```

Performance: $n = 5$

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
timeit.timeit(  
    'for x in xrange(20): tests.power_5(x)',  
    'import tests',  
    number=5000000)
```

statement	CPython 2.7	PyPy 1.9
power(5, x)	266.4 s	
power_5(x)	32.1 s	
x ** 5	12.9 s	

Performance: $n = 5$

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
timeit.timeit(  
    'for x in xrange(20): tests.power_5(x)',  
    'import tests',  
    number=5000000)
```

statement	CPython 2.7	PyPy 1.9
power(5, x)	266.4 s	0.73 s
power_5(x)	32.1 s	0.68 s
$x ** 5$	12.9 s	0.55 s

Performance: $n = 27$

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
timeit.timeit(  
    'for x in xrange(20): tests.power_27(x)',  
    'import tests',  
    number=500000)
```

statement	CPython 2.7	PyPy 1.9
power(27, x)	29.8 s	2.26 s
power_27(x)	6.65 s	2.18 s
$x ** 27$	15.8 s	21.9 s

Analysis: power(x, 27)

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

def power(x, n):	11.19 %
if not isinstance(n, int) or n < 0:	29.52 %
raise ValueError	
elif n == 0:	
return 1	
elif n % 2 == 0:	20.59 %
v = power(x, n / 2)	14.70 %
return v * v	5.80 %
else:	
return x * power(x, n - 1)	17.21 %

Analysis: power_27(x)

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

def power_27(x):	11.66 %
_pow_2 = x * x	9.55 %
_pow_3 = _pow_2 * x	5.58 %
_pow_6 = _pow_3 * _pow_3	6.58 %
_pow_12 = _pow_6 * _pow_6	6.95 %
_pow_13 = _pow_12 * x	6.54 %
_pow_26 = _pow_13 * _pow_13	17.93 %
_pow_27 = _pow_26 * x	14.58 %
return _pow_27	

Implementation

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def make_power(n):  
    ...  
    elif n == 0: return lambda x: 1  
    elif n == 1: return lambda x: x  
    elif n >= 2:  
        source = 'def fun(x):\n'  
        source += '\n'.join('    ' + s  
                             for s in _power_stmts(n))  
        fn_def = compile(  
            source, '<nofile>', 'exec')  
        eval(fn_def)  
        return locals()['fun']
```

Implementation

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def _power_stmts(n, ret=True):
    if n == 2: stmts = ['_pow_2 = x * x']
    elif n % 2 == 0:
        stmts = _power_stmts(n / 2, ret=False)
        stmts.append(
            '_pow_{n} = _pow_{n2} * _pow_{n2}'\
            .format(n = n, n2 = n / 2))
    else:
        stmts = _power_stmts(n - 1, ret=False)
        stmts.append('_pow_{n} = _pow_{n1} * x'\
            .format(n = n, n1 = n - 1))
    if ret: stmts.append('return _pow_{n}'.format(n=n))
    return stmts
```

To sum it up

- **Part 1: What is partial evaluation?**
- Part 2: When should we apply it?
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ We generate code, specializing on inputs
- ▶ Specialized code can run much faster
- ▶ But code generation is messy

Part 2

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

When should we apply it?

When should we apply it?

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ We can separate input into static and dynamic
- ▶ There is no single performance bottleneck
- ▶ The program is rather large and non-trivial
- ▶ So we don't want to rewrite it in C
- ▶ But we care about performance enough to introduce some complexity

Larger context

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ „Premature optimization is the root of all evil...”
- ▶ Profile
- ▶ Think about the algorithm
- ▶ Test coverage
- ▶ Benchmarks

Possible applications

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ Interpreters of:
 - ▶ programming languages
 - ▶ business-logic rules
 - ▶ DB queries
 - ▶ spreadsheet formulas
- ▶ Template engines

Pick low-hanging fruit first

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ Caching, memoization
- ▶ Closures
- ▶ Lazy evaluation

Example: closures

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def make_fn(static_input):  
    value = some_big_computation(static_input)  
    def fn(dynamic_input):  
        x = value.do_stuff(dynamic_input)  
        return x
```

But

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

If you need to change control flow, these methods will not help you.

No way to turn this

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power(x, n):  
    if not isinstance(n, int) or n < 0:  
        raise ValueError  
    elif n == 0:  
        return 1  
    elif n % 2 == 0:  
        v = power(x, n / 2)  
        return v * v  
    else:  
        return x * power(x, n - 1)
```

Into this

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

```
def power_27(x):  
    _pow_2 = x * x  
    _pow_3 = _pow_2 * x  
    _pow_6 = _pow_3 * _pow_3  
    _pow_12 = _pow_6 * _pow_6  
    _pow_13 = _pow_12 * x  
    _pow_26 = _pow_13 * _pow_13  
    _pow_27 = _pow_26 * x  
    return _pow_27
```

To sum it up

- Part 1: What is partial evaluation?
- **Part 2: When should we apply it?**
- Part 3: Real-world case study
- Part 4: Partial evaluation as a library

- ▶ Good fit for interpreters, but also other domains
- ▶ Test, benchmark, keep larger context
- ▶ Consider caching and closures

Part 3

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

Real-world case study

Real-world case study

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

- ▶ Problem domain: data collection, OLAP
- ▶ Data is stored in a graph DB, in EAV model
- ▶ Users define data structure, reports, formulas
- ▶ We must optimize formulas evaluation

The problem

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

- ▶ Formulas evaluation is the (big) bottleneck
- ▶ All simple and local optimizations are already done
- ▶ Time is evenly spread around a large set of functions
- ▶ And most of it is spent analysing formulas

Pseudocode

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

```
def get_value(attr, key):  
    if self.key_in_db(key):  
        return self.get_db_value(key)  
    for f in attr.formulae_before_aggregation():  
        res = self.get_formula_value(attr, key, f)  
        if res is not None:  
            return  
    for aggr_formula in attr.aggr_formulae():  
        ...  
    for f in attr.formulae_after_aggregation():  
        ...
```

Pseudocode

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

```
def get_formula_value(descriptor, key, formula):  
    lhs_key = dict(key)  
    if formula.left_filters:  
        cond_d, _ = formula.left_filters.as_dnf()  
        for k, v in cond_d:  
            if k in lhs_key:  
                ...  
    for arg in formula.args:  
        arg_key = dict(lhs_key)  
        if arg.filters:  
            ...  
        arg_value = get_value(arg.descriptor, arg_key)  
    ...
```

Gets compiled into this

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

```
def get_value_1(x, y, z):  
    ...  
    first_iter = True  
    while first_iter:  
        first_iter = False  
        arg_1 = get_value_2(x, z.next())  
        if arg_1 is None:  
            break  
    ...  
    res = some_fn(arg_1, arg_2)  
    ...
```


Results

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

- ▶ Calculation engine: 2000 LOC
- ▶ ``Interpreter" code that was specialized: 300 LOC
- ▶ ``Compiler" code, that does this specialization: 400 LOC
- ▶ Performance: about 10x faster

Performance

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library



Making it maintainable

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

- ▶ Good test coverage
- ▶ Benchmarks, regression tests on speed and correctness
- ▶ Debugging: output generated code to file, support pdb
- ▶ Can run without generating specialized code
- ▶ New features are added first without specialization

To sum it up

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- **Part 3: Real-world case study**
- Part 4: Partial evaluation as a library

- ▶ Partial evaluation gave 10x speedup
- ▶ Removed a huge bottleneck
- ▶ Code it still maintainable
- ▶ But you have to be careful

Part 4

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

Partial evaluation as a library

Partial evaluation as a library

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

How cool would it be to just write:

```
import ast_pe
power_5 = ast_pe.specialized_fn(
    power, n=5)
```

Partial evaluation libraries

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ Were developed for Lisp dialects, and even for C
- ▶ Was attempted for Ruby
- ▶ I have not found a Python one
- ▶ So decided to write it

How does it work

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ Transforms AST of the target function
- ▶ Looks like a little Python AST interpreter
- ▶ Uses Python interpreter to evaluate expressions depending only on static input

Abstract Syntax Tree

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

```
def sample_fn(x, y, foo='bar'):  
    if (foo == 'bar'):  
        return x + y  
    else:  
        return x - y
```

Abstract Syntax Tree

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

```
>> import ast, inspect, meta.asttools
>> meta.asttools.print_ast(
    ast.parse(inspect.getsource(sample_fn)))
```

Abstract Syntax Tree

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

```
FunctionDef(  
    args=arguments(  
        args=[Name('x'), Name('y'), Name('foo')],  
        defaults=[Str(s='bar')],  
        kwarg=None,  
        vararg=None),  
    body=[If(  
        body=[Return(value=BinOp(  
            left=Name(ctx=Load(), id='x'), op=Add(),  
            right=Name(ctx=Load(), id='y')))],  
        orelse=[Return(value=BinOp(  
            left=Name(ctx=Load(), id='x'), op=Sub(),  
            right=Name(ctx=Load(), id='y')))],  
        test=...
```

Transforming AST

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

```
class PartialEvaluator(ast.NodeTransformer):  
    ...  
    def visit_Name(self, node):  
        self.generic_visit(node)  
        if isinstance(node.ctx, ast.Load) \   
            and node.id in self.bindings:  
            value = self.bindings[node.id]  
            if isinstance(value, numbers.Number):  
                return ast.Num(n=value)  
        ...
```

But this is not so easy

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ I/O - files, DB: `f.write(data)`
- ▶ Variable assignment:
`static_input += dynamic_i`
- ▶ Variable mutation:
`static_input.append(dynamic_i)`
- ▶ Redefining method:
`some_obj.method = another_method`
- ▶ Redefining operator:
`def __getitem__(self, key, ...):`

So we have to

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ Do dataflow analysis
- ▶ Detect mutations and assignments
- ▶ Rollback if some assumptions were proven wrong

Assumptions

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ You do not mutate/change static input
- ▶ Mark pure functions (`@pure` decorator)
- ▶ Mark functions you want to be inlined (`@inline` decorator)

To sum it up

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ It is mostly a prototype now
- ▶ But proves it can be done

Contacts

- Part 1: What is partial evaluation?
- Part 2: When should we apply it?
- Part 3: Real-world case study
- **Part 4: Partial evaluation as a library**

- ▶ Konstantin Lopuhin
- ▶ `konstantin.lopuhin@chtd.ru`
- ▶ `https://github.com/chtd/ast_pe`
- ▶ CHTD - `http://chtd.ru`
- ▶ Twitter: `@python_chtd`